Anvit More

ML Systems · GPU Kernels · LLM Inference · Applied RL | anviit22@gmail.com | linkedin.com/in/anvit22 | github.com/anviit | Portfolio | +91 9921243173 | Pune, India | Open to relocate: DE / CH / EU | Requires work permit

PROFILE

ML Systems Engineer specialising in Triton/CUDA kernel engineering, production LLM inference, and applied RL — familiar with vLLM, SGLang, and TGI, with hands-on understanding of the KV cache and memory allocation problems they solve. Production background in medtech AI at Plus91: clinical decision support, OCR, ASR, and radiology tooling across pilot and live deployments. MSc Data Science, University of Edinburgh (B+). German B1. Open to ML Systems, LLM Inference, and Automotive AI roles across Germany, Switzerland, and the EU.

EXPERIENCE

ML & AI Engineer · Plus91 Technology Pvt. Ltd., Pune *Mar 2025 – Present*

- Clinical Decision Support: Fine-tuned Mistral 7B v0.3 on USMLE MCQs, clinical textbooks, and hospital data — 10 ranked differentials with reasoning chains. 90%+ diagnostic accuracy in clinician review. Active pilot with clinical validation ongoing.
- Intelligent OCR Pipeline: Vision-LLM (OpenAI/Gemini Vision + Vertex AI fallback) processing 200-500 TRFs/day at 90%+ accuracy. Event-driven AWS S3 routing (failure/success buckets). Turnaround ~15 min → under 20 sec (~85% reduction).
- Audio Scribe: Speech-to-structured-note for 60+ min consultations (General, Oncology, Ophthalmology). Audio served via AWS S3. WER <10%, 500+ consultations in production.
- Radiology Dictation Engine: Voice-driven report editor with BI-RADS template support. Reduces report editing to under 20 sec. In production rollout.
- 3rd Place / 100+ Teams — NHA NHCX Hackathon 2025: PDF-to-FHIR R4 microservice, 13+ resource types, 4-tier offline validator. [Certificate]

SELECTED PROJECTS

LLM Training Kernels — Triton from Scratch · github.com/anviit/triton-llm-kernels

- FlashAttention 2.52x, Bias+GELU 14.65x (172 GB/s), AdamW 3.45x (177 GB/s), inference attention 3.94x. Validated against fp32 reference.

Phi-3 Mini Decode Engine · github.com/anviit/phi3-triton-decode

- Custom Triton GQA decode kernel. 2.5x speedup (307ms vs 775ms/token). 0.65% of CUDA time is attention — gain from eliminating torch.cat KV reallocation across 32 layers. Same insight as vLLM/SGLang/TRT-LLM.

Production LLM Inference Stack · github.com/anviit/llm-inference-serving

- TinyLlama 1.1B on 6 GB GPU: 28ms TTFT P50, 39.4 tok/s (~85% mem BW), 81% Redis cache hit, 100% success at concurrency=10.

Autonomous Race Strategy Engine · github.com/anviit/autonomous-strategy-engine

- PPO agent +58% points vs rule-based (E[Pos] 1.00 vs 3.09). Monte Carlo planner at 870 rollouts/sec. Directly maps to ADAS planning and EV energy management.

EDUCATION

MSc Data Science · University of Edinburgh, UK *2023 – 2024*

Grade B+ · Dissertation: RL-based Hospital Resource Management (Ophthalmology) · Applied ML, Network Science, Bayesian Methods, Probabilistic ML

BSc Computer Science · Fergusson College, Pune, India *2020 – 2023*

Grade O (Outstanding) · 9.6 CGPA · Final Project: Decision Engine for Automobile Sector — user requirements to cost/benefit optimisation

TECHNICAL SKILLS

GPU / Kernels: Triton, CUDA, PyTorch FP16, triton.testing.do_bench, CUDA Events, pynvml
LLM / Inference: PyTorch, HF Transformers, FastAPI, Redis, uvicorn, Whisper ASR, Gemini Vision, LoRA/QLoRA, FAISS | Familiar: vLLM, SGLang, TGI
RL / Simulation: Stable-Baselines3, Gymnasium, NumPy, SciPy
Architectures: FlashAttention, KV Cache, INT8 Quantization, PPO, Monte Carlo, PEFT, GQA, RoPE
MLOps / Cloud: Docker, Git, pytest, Bash, AWS S3, GCP Vertex AI
Languages: English C2 / Native · German B1 · Spanish A1